

**Name: Krishna Somani**

**Roll No. : 2021058**

## **Report - Assignment 1**

### **3. Query Implementation using Apache Spark**

- 1) Fetching all students enrolled in a specific course.

+---+-----+-----+  _id first_name last_name  +---+-----+-----+		
10	Dre	Kersley
1	Sharyl	Jaouen
11	Gabi	Crystal
61	Victor	Sleany
56	Tome	Maven
13	Jo-ann	Heindrich
57	Nikkie	Eadmeades
16	Collette	Brenneke
+---+-----+-----+		

- 2) Calculating the average number of students enrolled in courses offered by a particular instructor at the university.

+-----+  average_students  +-----+	
	3.0
+-----+	

- 3) Listing all courses offered by a specific department.

_id	course_name	course_code
8	Calculus	MATH101
2	Linear Algebra	MATH201
14	Statistics	MATH301
20	Discrete Mathematics	MATH401
26	Differential Equa...	MATH501
32	Number Theory	MATH601
38	Topology	MATH701
44	Abstract Algebra	MATH801
50	Real Analysis	MATH901
56	Number Theory	MATH1001

- 4) Finding the total number of students per department.

department_name	total_students
Art	51
Chemistry	53
English	49
History	42
Music	49
Mathematics	48
Physics	47
Computer Science	65
Economics	39
Biology	39

- 5) Finding instructors who have taught all the BTech CSE core courses sometime during their tenure at the university.

instructor_id	instructor_name
1	Cariotta Worboy

- 6) Finding top-10 courses with the highest enrollments.

course_id	course_name	total_enrollments
3	Classical Mechanics	17
8	Calculus	9
5	Genetics	8
1	Introduction to C...	8
9	Quantum Mechanics	6
40	Environmental Che...	6
6	Data Structures	5
7	Artificial Intell...	4
2	Linear Algebra	4
4	Organic Chemistry	3

## 4. Optimizations

### Techniques Used:

#### 1. Indexing

```
db.instructors.createIndex({"instructor_id": 1 })
db.students.createIndex({"student_id": 1 })
db.courses.createIndex({"course_id": 1 })
db.courses.createIndex({"course_id": 1, "instructor_id": 1 })
```

Indexes in MongoDB enable fast document retrieval by allowing quick access to relevant records without full collection scans. Indexing fields like `student\_id`, `course\_id`, and `instructor\_id` decreases search times and input/output operations, enhancing performance. Compound indexes (e.g., on `course\_id` and `instructor\_id`) further improve efficiency for multi-field queries, reducing post-retrieval filtering. This is vital for maintaining consistent performance in read-heavy applications as the database scales

#### 2. Partitioning (In Spark)

On basis of columns:

```
courses_partitioned_df = courses_df.repartition(n, "_id")
students_df = students_df.repartition(n, "_id")
instructors_partitioned_df = instructors_df.repartition(n, "_id")
```

S. No.	Exec_time_before_optimization	Exec_time_after_optimization
1	0.3738 seconds	0.2509 seconds
2	0.9351 seconds	0.4811seconds
3	0.1162 seconds	0.1156 seconds
4	0.3713 seconds	0.309 seconds
5	0.5245 seconds	0.2633seconds
6	0.3862 seconds	0.3501 seconds

It can be seen these optimizations lead to some decrease in execution time, but they can significantly perform well particularly when used with larger datasets or in distributed computing settings. However, with smaller datasets, the overhead associated with creating indexes or repartitioning may occasionally cause longer query execution times. So, overall results don't show too much of a difference!